

Vision par Ordinateur : Indexation et recherche d'images

Database indexing

Thanks to Laurent Amsaleg and Hervé Jégou for some of this material

Context: Similarity of Images

- Comparing description data (feature vectors) instead of images directly
 - Typically high-dimensional vectors
 - Vectors define points in high-dim spaces
- The similarity between images is in proportion to the similarity of their feature vectors
- Two images are said to be similar if their descriptors are close in the high-dim space
- Everything relies on a metrics between vectors
 - often a distance or a similarity
 - all dimensions involved

Context: Image retrieval

- Goal:
 - Given a feature vectors database,
 - Given a query image described by a set of query vectors
- ⇒ retrieve the closest feature vectors of the database:
k-ppv or ϵ -sphere
- ⇒ return a **ranked** list of images:
Voting mechanism

3

Two types of Searches

- Type 1 : **K-nn search**
 - searching for K nearest neighbors
 - result set of fixed size
 - near does not means close
- Type 2 : **ϵ search**
 - searching within a "ball":
 $\|p - p_i\| \leq \epsilon$
 - bounds dissimilarity
 - result set of unpredictable size



Toy Example

- 1 query image
- Database: 2 images I_1 et I_2



- Description:
 - local descriptors of dimension 3
 - Distance L_1

5

Toy Example

- Number of distances to compute?
- Ranked list of similar images?

- Query descriptors:

	x1	x2	x3
q1	5	5	1
q2	3	6	1
q3	5	4	6
q4	5	1	7


- DB descriptors:

	x1	x2	x3
d11	4	2	1
d12	0	7	7
d13	6	3	5

	x1	x2	x3
d21	5	0	3
d22	5	3	4
d23	6	5	2
d24	0	1	7
d25	4	7	1
d26	3	2	0

6

Toy Example



7

Toy Example

Computation of 36 distances L_1 between descriptors:

	q1	q2	q3	q4
d11	4	5	8	8
d12	13	10	9	11
d13	7	10	3	5
d21	7	10	7	5
d22	5	8	3	5
d23	2	5	6	10
d24	15	14	9	5
d25	3	2	9	13
d26	6	5	10	10

For each query vector:

⇒ retrieve the closest feature vectors of the database:

k-ppv or ϵ -sphere

⇒ return then a **ranked** list of images:

Voting mechanism

8

Naïve, straightforward approach

- [Compute all the descriptors]
- Descriptors are all stored in a file
 - one after the other, sequentially
- Search Process:
 - read the file (large chunks at once)
 - compute distance between query and descriptors coming from disk
 - keep track of the K nearest neighbors (for ex)
- Exhaustive search, sequential

Naïve, straightforward approach

- Piece of cake!
- But very costly:
 - fetch data from disks (lots of I/Os)
 - compute distances over ALL data
- Not very realistic when dealing with:
 - very large volume of data
 - high dimensionality
 - complex metrics (EMD)

Python time: exhaustive search

- Generate a random matrix of 10 descriptors of dimension 5
- Given a query vector:
 - Use FLANN to retrieve the 3 k-nn
- Given several query descriptors:
 - Use FLANN to retrieve the 3 k-nn
- Print execution time to retrieve the 50-nn for d=500, and:
 - 1 query, N=10,000
 - 1 query, N=100,000
 - 1 query, N=1,000,000
 - 100 queries, N=10,000
 - 100 queries, N=1,000,000

11

How can we accelerate the search?

12

We need...

- To reduce the volume of data to analyze during searches
- To enclose searches
- To efficiently access data and return answer
- How? By structuring the descriptors

*Doing this fast
requires multidimensional indexing*

How Can We Index Data?

- Enclose the search
 - descriptors live in a high-dimensional space
 - analyze only interesting regions of space
- Key Idea:
 - group descriptors into **cells**
 - ~ much fewer cells than descriptors
 - detect useless cells and ignore them
 - ~ the ones containing descriptors that can not be part of the final result
 - fine-grain analysis of remaining cells
 - ~ compute distances using the descriptors they contain


Table of Contents

1. Data indexing: why we need cells?
2. How to construct cells?
3. High-dimensional weirdness
4. Approximate searches

Why Relying on Cells?

- First, compute distances from query to cells
- Many cells can be ignored without accessing any vector they contain
 - much fewer cells than descriptors
 - low complexity
 - simply need geometrical information on cells
 - great response time gains

Geometry Helps




- Here is a data collection
- Searching for the 4 NN of a query point

Exercise Time...


- Goal:
 - find simple geometric rules allowing to decide whether the contents of a particular cell might or might not be useful for answering a particular query
- Hints:
 - it's based on minimum and maximum distances from the query point to cell boundaries

Ready?




- Searching for the 4 NN of q
- blue cell obviously sufficient
- How can we eliminate red cell?

Go!



Ta daaa...


answer is qB & qC



- can reject red without analyzing its contents

Your turn: more difficult (?)

- What happens in this case?
 - ~ still searching for the 4 NN of q




Question 2

- Can we prune green?
- Can we prune purple?
- Can we prune both?
- Can we prune none?

Ta da daaaa...

- Searching for the 4 NN of q



Purple is really close, so analysis needed

cells overlap

Can not reject green. Its contents needs analysis

Answer= can prune none

Geometry Helps

- Compute distance to all cells
- if $d_{\min}(q, C_i) \geq d_{\max}(q, C_i)$ then reject C_i
 - few cells: few distance calculations
 - very few candidates: few distance calculations
 - great gain in performance!
 - make sure you have k nn

Geometry Helps Again!

- Previous rule computes a list of candidate cells
- List ordered w.r.t. increasing distance to q
 - fetch first cell
- Process all vectors in the current cell
 - update list of nearest neighbors
- if $d_{\min}(q, C_i) \geq d(q, nn_k)$ then stop search
 - else fetch next cell; loop

Indexing =



- Design algorithms for building cells
- Design data structures to organize cells
- Design algorithms traversing data structures efficiently

Table of Contents


1. Data indexing: why we need cells?
2. How to construct cells?
3. High-dimensional weirdness
4. Approximate searches

Grouping Descriptors into Cells?


- Two broad approaches for cell construction
 - based on distances between descriptors
 - based on high-dim space partitioning
- Information stored with cells
 - position in space, size, centre, population, ...
 - Usually: minimum bounding rectangle/sphere



example : 2d feature space




Grouping according to the partitioning of the feature space




Grouping according to data proximity

often rectangle or spheres
often minimum bounding



Multidimensional Indices


- Traditional indices for DBMS: trees




- Many indexing schemes have been designed to handle multidimensional data
- Overview of the 2 seminal approaches
 - R-Tree: data proximity
 - KD-Tree: space partitioning

R-Tree

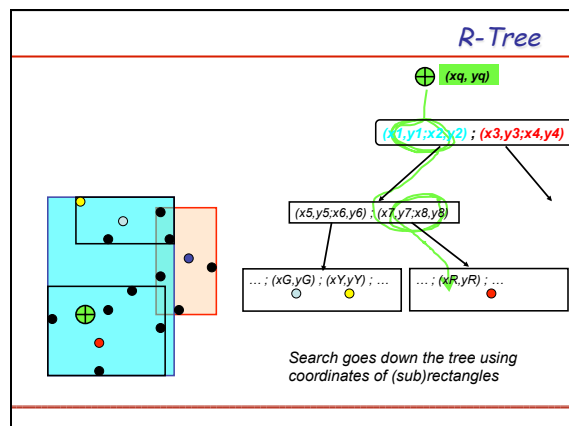
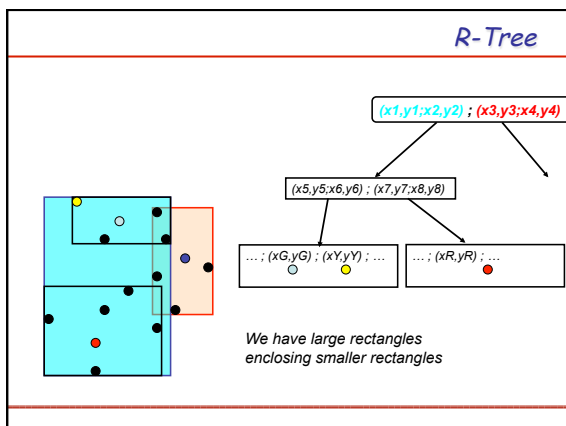




- Guttman, 1984, *ACM SIGMOD*
- B-Tree extended to high-dim spaces
 - nodes & leaves = hyper-rectangles
- Dynamic insertion of vectors
- Leaves (fixed size) group vectors
- Leaves split on overflow
 - update parent
- Balanced tree



$(a1, b1) ; (c1, d1)$




$(a1, b1) ; (c2, d2)$



Building One R-Tree

- Nodes and leaves have a fixed size
 - leaves are on disk
 - often a multiple of I/O granule (128kb)
- Nodes and leaves must split on overflow
 - many options for splitting
 - always problematic: causes overlap




R-Tree

- Minimum Bounding Rectangles
- Overlapping exists
- Tree structure makes traversal fast (log)
- At the roots of almost all indices having cells based on the proximity of their points
 - SS-Tree: hyperspheres
 - SR-Tree: hyperrectangles \cap hyperspheres
 - ...

Kd-Tree


- Bentley, 1975. Communication of the ACM.
- This is a space partitioning strategy
- Space split using hyperplanes
 - perpendicular to one axis
 - splitting point = median of points on that axis
 - pick next axis

Kd-Tree




Kd-Tree


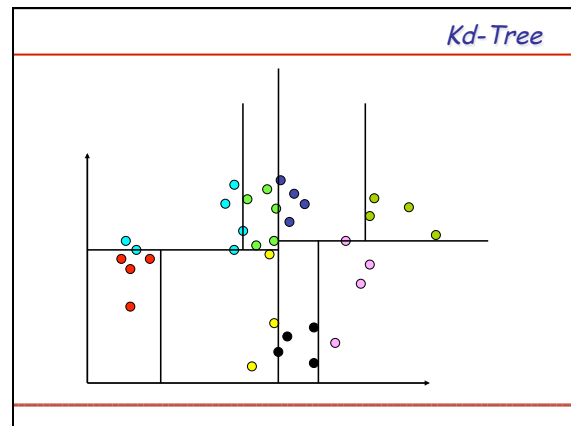
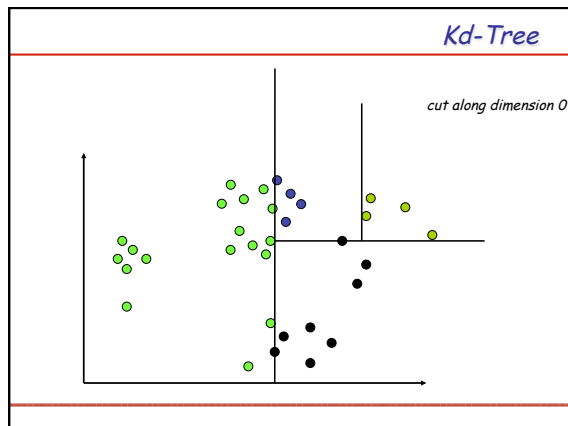
cut along dimension 0



Kd-Tree

cut along dimension 1





- ### Kd-Tree
- All cells are disjoint
 - no overlap between cells
 - super-cool for pruning rules!
 - Nodes in memory
 - Leaves on disk
 - At the roots of many space-partitioning approaches
 - BSP Tree (Binary Space Partitioning)
 - ~ subdivision lines are not parallel to coordinate system
 - ~ this leads to Voronoi cells
 - Variations
 - ~ where to split (often median)
 - ~ what dimension to use next
 - ~ balanced tree, or not...

- ### Python time: indexing the DB
- Compare the execution times to retrieve the 50-nn for $d=500$, and:
 - 1000 queries, $N=1,000,000$
- When using a brute-force approach (linear search) and a KD-Tree
- 53

- ### Are these good approaches?
- What about their:
 - pruning power
 - speedup w.r.t. sequential scan
 - construction complexity
 - ...
 - Do experiments and/or study high-dimensional spaces

Table of Contents

1. Data indexing: why we need cells?
2. How to construct cells?
3. High-dimensional weirdness
4. Approximate searches

Properties of High-Dimensional Spaces

- Descriptors are typically high-dimensional
- Some weird things happen in such spaces
 - algo fine at low-dim die at high-dim
- Watch your intuition. Things in 2d or 3d are way different
- This is called **curse of dimensionality**

Python time: high-dimensional spaces

- Study a uniform **IID** distribution, when d grows
 - Generate a DB containing N vectors of dimension d ($N, d = \text{parameters}$), uniformly distributed
 - Generate one or several query vector(s) of dimension d
 - For $d=2$, for each query, compute:
 - ~ The distance its nearest neighbor (1-nn)
 - ~ The distance to the farthest neighbor (N-nn)
 - ~ The ratio between the 1-nn and 2-nn (Low matching criteria)
 - ~ The ratio between the 1-nn and N-nn
 - Repeat for increasing values of d : 3, 5, 10, 50, 100, 200, 500, 1000
- What happens?

57

1) Vanishing Variance

- Shaft [ACM TODS 2006] have observed that, with an **IID** distribution, when d grows:
 - the distance between all pairs of points tend to be similar
 - both $\text{dist}(q, \text{nn})$ and $\text{dist}(q, \text{farthest neighbor})$ grow
 - $\text{dist}(q, \text{nn})$ grows faster than $\text{dist}(q, \text{fn})$
 - the closest point to a given one is as far as its farthest
 - NN becomes very unstable
 - this is called "Vanishing Variance"
- Some data sets can hardly be indexed

2) High-dim Space Partitioning

- The number of cells for space partitioning grows exponentially with d
- Example:
 - $d = 30$: split each dimension in 2
 - $2^{30} = 1\,073\,741\,824$ cells
 - much more cells than data points
 - many empty cells, costly to keep track of them
 - unlikely to have more than 1 point per cell
- Partitioning for cells is useless


2) High-dim Space Partitioning

- If a point is near a frontier, then must analyze the neighboring cell
- How many cells touch any particular cell when $d=30$ and each split in 2?
- How likely is it that a random point gets near a frontier?
- Example: In $[0, 1]^d$, what is the probability for a point (uniformly sampled) to be at less than δ from the frontier?
- Python: compute this proba for $d=1, 2, 3, 5, 10, 50, 100, 500$

2) High-dim Space Partitioning

- In $[0,1]^d$, the probability to be at less than δ from the frontier is:

$$P = 1 - (1 - 2\delta)^d$$



- δ is small here (1%)

2) High-dim Space Partitioning

- All data points are near frontiers
 - must check on the other side
- One cell has very many neighbors
 - must check many neighboring cells
- Many cells are empty
- Vectors are lonely in cells
 - cells are useless
- Nice, isn't it?

3) Rectangles and Spheres

- Simple shapes
 - easy to encode, cheap for geometrical rules
- Let's compare them in high-dim, $d=30$
- Compactness
 - volume = 1 \Rightarrow diagonal = 5.47, radius = 1.43
 - hyperspheres are much more compact than hyperrectangles are (corners take up space)

3) Rectangles and Spheres


- Enlargement
 - to absorb a new data point, shape must expand
 - diagonal ± 0.5 ; radius ± 0.5
 \Rightarrow volume R=35; volume S=126
 - hyperrectangles grow much slower than hyperspheres do

3) Rectangles and Spheres

- Consequences
 - in high-dim, rectangles are huge
 - in high-dim, spheres get fat
 - there is a lot of overlap
 - the more overlap, the less geometrical rules can be effective
 - so the more cells you need to analyze
- Cells are useless
 - isn't this cool?

4) Rectangles and Spheres

- Let's zoom on the volume of R or S



- this is the shell of the shape


4) Rectangles and Spheres

- the volume of the shell grows exponentially with d
- when d large, all the volume is in the shell
 - eating high-dim eggs will not stuff you
- it is unlikely anything can be in the "center"
- so again, frontier and enlargement problems
- Isn't creating cells a stupid idea?

5) Empty Space Phenomenon

- Whatever you do, when d grows, the space becomes more and more empty
- Example:
 - S is an hypersphere with radius of 0.5 in $[0,1]^d$
 - compute N to have at least 1 vector in S

$d=2$



$$V_d(r) = \frac{\pi^{d/2} r^d}{\Gamma(1+d/2)}$$

when d even : $\Gamma(1+d/2) = (d/2)!$

when d odd : $\Gamma(1+d/2) = \sqrt{\pi} \frac{d!!}{2^{(d+1)/2}}$


with $d!! = \begin{cases} 1, & d = -1, 0, 1 \\ d(d-2)!!, & d \geq 2 \end{cases}$

5) Empty Space Phenomenon

- Let's increase d
 - $d=10$, $N \approx 400$
 - $d=20$, $N \approx 40 \cdot 10^6$

$d=40$, $N=3000 \cdot 10^{20}$
 $d=100$, $N=5000 \cdot 10^{69}$

- Why is that?



- "Corners" are consuming all the space

Indexing High-dim Data

- If cells are overlapping
- If most of cells are empty
- If there is only one point per cell
- If you need to analyzed all neighboring cells

- Then:

Read all cells, randomly
 => Lot of (random) I/O

=> Much more costly than a sequential scan

Indexing High-dim Data

- This is very true for uniform data
 - theoretically proven - what can you do...
 - no indexing method will beat seq scan
 - particularly when searching exact neighbors
- Less true when data is not uniform
 - real data is more contrasted
 - vanishing variance not kicking so strongly
 - but high dimensionality still causes trouble

Summary

- Consensus: only the sequential scan has a fair behavior
- All other approaches eventually fail, and degenerate to worse.
- What the hell is going on during a search?


High-Dim. Data - What is going on?

- What are the effects of all this?
- What is happening to search processes?
- Elements for understanding:
 - years of experiments and observations
 - results extracted from a paper: "A Case Study of the Quality vs. Time Trade-off for Approximate Image Descriptor Search", Sigurðardóttir, Hauksson, Jónsson, Amsaleg, EMMA'05 (with ICDE)


Experiments

- 5,017,298 descriptors, 24 dimensions
- 52,273 real-life images (news, photo agency)
- 2 cell forming approaches
 - SR-Tree
 - BAG

of chunks needed to get NN



Time to get NN



Overall

- You very quickly get the first NN
 - first cell(s) very profitable
- You get others more slowly
 - next cells less and less profitable
 - many NN \Rightarrow many cells
- Termination problem
 - overlap \Rightarrow many cells candidate (all!)
 - no way to say STOP

The Coolest Idea

- Why waiting 10h to get a perfect result if in 10 sec we can get a pretty good one?
- Approximate Searches
 - trade response time against result quality
 - this is where the action is!